

Project Track 1 Stage 3

Database Implementation	2
GCP Database Tables - "Screenshot of connection (terminal/command-line information)"	2
DDL commands for tables	2
Screenshot of count query for proof of 1000 rows in three different tables	3
Advanced Queries + Indexing Analysis	4
Query 1: Most Popular Products (Products found most frequently in user bags)	4
Indexing Design 1	4
Indexing Design 2	5
Indexing Design 3	6
Final Indexing Choice and Why	7
Query 2: User Bag Overview During Friends Search	8
Indexing Design 1	8
Indexing Design 2	9
Indexing Design 3	9
Final Indexing Choice and Why	9
Query 3: Users also bagged	11
Indexing Design 1	12
Indexing Design 2	12
Indexing Design 3	13
Final Indexing Choice and Why	13
Query 4: Aggregations for common products bagged between the application user and the user's bag they are viewing	13
Indexing Design 1	14
Indexing Design 2	15
Indexing Design 3	15
Final Indexing Choice and Why	15
Appendix	16

Database Implementation

GCP Database Tables - “Screenshot of connection (terminal/command-line information)”

```
CLOUD SHELL
Terminal (vanity-cs411-sp24) x + v

mysql>
mysql>
mysql>
mysql> show databases;
+-----+
| Database |
+-----+
| MojoDojoDB |
| information_schema |
| mysql |
| performance_schema |
| sys |
+-----+
5 rows in set (0.00 sec)

mysql> use MojoDojoDB;
Database changed
mysql> show tables;
+-----+
| Tables_in_MojoDojoDB |
+-----+
| BagItems |
| Brands |
| ProductTags |
| Products |
| Reviews |
| Tags |
| Users |
+-----+
7 rows in set (0.01 sec)

mysql> SELECT * FROM Products LIMIT 3;
+-----+
| ProductId | ProductName | Size | Price | LikeCount | BrandId |
+-----+
| P01018539 | Master Mattes Liquid Eyeliner | 0.04 oz/ 1.2 mL | 24 | 13625 | 6332 |
| P02318798 | Master Metallics Eyeshadow Palette | 12 x 0.03 oz/ 1 g | 50 | 25738 | 6332 |
| P04141879 | Liquid Touch Foundation Brush | | 29 | 98192 | 6317 |
+-----+
3 rows in set (0.01 sec)

mysql> 
```

DDL commands for tables

```
CREATE TABLE Users(UserId INT PRIMARY KEY, UserName VARCHAR(30) NOT NULL, Password VARCHAR(30) NOT NULL, FirstName VARCHAR(255) NOT NULL, LastName VARCHAR(255) NOT NULL, Email VARCHAR(255) NOT NULL);
```

```
CREATE TABLE Brands(BrandId INT PRIMARY KEY, BrandName VARCHAR(255) NOT NULL, BrandRating REAL);
```

```
CREATE TABLE Products(ProductId VARCHAR(64) PRIMARY KEY, ProductName VARCHAR(255) NOT NULL, Size VARCHAR(128), Price INT, LikeCount INT, BrandId INT NOT NULL, FOREIGN KEY (BrandId) REFERENCES Brands(BrandId));
```

```
CREATE TABLE Reviews(ReviewId INT PRIMARY KEY, Rating INT, Text
VARCHAR(1000), Title VARCHAR(255), Date DATE, ProductId VARCHAR(64) NOT
NULL, FOREIGN KEY (ProductId) REFERENCES Products(ProductId));
```

```
CREATE TABLE Tags(TagId INT PRIMARY KEY, TagName VARCHAR(64) NOT NULL,
Standing INT);
```

```
CREATE TABLE BagItems(UserId INT NOT NULL, ProductId VARCHAR(64) NOT
NULL, DateAdded DATE, PRIMARY KEY (UserId, ProductId), FOREIGN KEY
(UserId) REFERENCES Users(UserId), FOREIGN KEY (ProductId) REFERENCES
Products(ProductId));
```

```
CREATE TABLE ProductTags(ProductId VARCHAR(64) NOT NULL, TagId INT NOT
NULL, PRIMARY KEY (ProductId, TagId), FOREIGN KEY (ProductId)
REFERENCES Products(ProductId), FOREIGN KEY (TagId) REFERENCES
Tags(TagId));
```

```
CREATE TABLE ProductClusters(ProductId VARCHAR(64) NOT NULL, ClusterId
INT NOT NULL, PRIMARY KEY (ProductId, ClusterId), FOREIGN KEY
(ProductId) REFERENCES Products(ProductId));
```

Screenshot of count query for proof of 1000 rows in three different tables

Products

```
mysql> SELECT count(*) FROM Products;
+-----+
| count(*) |
+-----+
|      8495 |
+-----+
1 row in set (0.06 sec)
```

Reviews

```
mysql> SELECT count(*) from Reviews;
+-----+
| count(*) |
+-----+
|     17170 |
+-----+
1 row in set (0.34 sec)
```

Tags

```
mysql> SELECT count(*) FROM Tags;
+-----+
| count(*) |
+-----+
|      4329 |
+-----+
1 row in set (0.00 sec)

mysql> █
```

Advanced Queries + Indexing Analysis

Query 1: Most Popular Products (Products found most frequently in user bags that have been viewed in-app, have more than 1,000 likes, and have high review ratings)

```
SELECT DISTINCT Pro.ProductName, B.BrandName, Subquery.TimesBagged
FROM
(SELECT Bag.ProductId, COUNT(Bag.ProductId) as TimesBagged FROM BagItems Bag
GROUP BY Bag.ProductId) Subquery
LEFT OUTER JOIN Products Pro ON Pro.ProductId = Subquery.ProductId
LEFT OUTER JOIN Brands B ON B.BrandId = Pro.BrandId
LEFT OUTER JOIN Reviews R on R.ProductId = Pro.ProductId

WHERE (Subquery.TimesBagged >= 0.8 * (SELECT MAX(Subquery2.TimesBagged) FROM
(SELECT Bag.ProductId, COUNT(Bag.ProductId) as TimesBagged FROM BagItems Bag
GROUP BY Bag.ProductId) Subquery2)) AND Pro.LikeCount > 1000 AND Pro.ViewCount
> 1 AND R.Rating > 4;
```

Output: Only 3 Most Popular Products in User bags (this table grows with user activity)

```
+-----+-----+-----+
| ProductName | BrandName | TimesBagged |
+-----+-----+-----+
| The Camellia Oil 2-in-1 Makeup Remover & Cleanser | Tatcha | 20 |
| Mini Superfood Antioxidant Cleanser | Youth To The People | 20 |
| Mini Revitalizing Supreme+ Youth Power Creme Moisturizer | Estée Lauder | 19 |
+-----+-----+-----+
3 rows in set (0.01 sec)
```

EXPLAIN ANALYZE Performance before new indexes:

```
-----+
| -> Table scan on <temporary> (cost=68.37..70.67 rows=9) (actual time=0.601..0.601 rows=3 loops=1)
|   -> Temporary table with deduplication (cost=68.07..68.07 rows=9) (actual time=0.600..0.600 rows=3 loops=1)
|     -> Nested loop inner join (cost=67.20 rows=9) (actual time=0.445..0.574 rows=3 loops=1)
|       -> Nested loop left join (cost=43.90 rows=9) (actual time=0.380..0.428 rows=5 loops=1)
|         -> Nested loop inner join (cost=40.85 rows=9) (actual time=0.372..0.410 rows=5 loops=1)
|           -> Filter: (Subquery2.TimesBagged >= <cache>((0.8 * (select #3)))) (cost=70.32..13.32 rows=79) (actual time=0.349..0.364 rows=6 loops=1)
|             -> Table scan on Subquery (cost=71.07..76.50 rows=236) (actual time=0.178..0.190 rows=102 loops=1)
|               -> Materialize (cost=71.05..71.05 rows=236) (actual time=0.176..0.176 rows=102 loops=1)
|                 -> Group aggregate: count(Bag.ProductId) (cost=47.45 rows=236) (actual time=0.061..0.144 rows=102 loops=1)
|                   -> Covering index scan on Bag using idx_bagitems_productid (cost=23.85 rows=236) (actual time=0.054..0.098 rows=237 loops=1)
|                     -> Select #3 (subquery in condition; run only once)
|                       -> Aggregate: max(Subquery2.TimesBagged) (cost=100.10..100.10 rows=1) (actual time=0.151..0.151 rows=1 loops=1)
|                         -> Table scan on Subquery2 (cost=71.07..76.50 rows=236) (actual time=0.130..0.141 rows=102 loops=1)
|                           -> Materialize (cost=71.05..71.05 rows=236) (actual time=0.129..0.129 rows=102 loops=1)
|                             -> Group aggregate: count(Bag.ProductId) (cost=47.45 rows=236) (actual time=0.026..0.106 rows=102 loops=1)
|                               -> Covering index scan on Bag using idx_bagitems_productid (cost=23.85 rows=236) (actual time=0.024..0.064 rows=237 loops=1)
|                                 -> Filter: ((Pro.LikeCount > 1000) and (Pro.ViewCount > 1)) (cost=0.25 rows=0.1) (actual time=0.007..0.007 rows=1 loops=6)
|                                   -> Single-row index lookup on Pro using PRIMARY (ProductId=Subquery.ProductId) (cost=0.25 rows=1) (actual time=0.007..0.007 rows=1 loops=6)
|                                     -> Single-row index lookup on B using PRIMARY (BrandId=Pro.BrandId) (cost=0.26 rows=1) (actual time=0.003..0.003 rows=1 loops=5)
|                                       -> Limit: 1 row(s) (cost=2.40 rows=1) (actual time=0.028..0.028 rows=1 loops=5)
|                                         -> Filter: (R.Rating > 4) (cost=2.40 rows=3) (actual time=0.028..0.028 rows=1 loops=5)
|                                           -> Index lookup on R using idx_Reviews_ProductId_ReviewId (ProductId=Subquery.ProductId) (cost=2.40 rows=9) (actual time=0.027..0.028 rows=1 loops=5)
|
|-----+
```

Total Cost: 70.67 + 68.07 + 67.20 + 43.90 + 40.85 + 70.32 + 76.50 + 71.05 + 47.45 + 23.85 + 100.10 + 76.50 + 71.05 + 47.45 + 23.85 + 0.25 + 0.25 + 0.26 + 2.40 + 2.40 + 2.40 = **906.77**

Indexing Design 1

- New Index:
CREATE INDEX idx_Products_LikeCount ON Products(LikeCount);
- EXPLAIN ANALYZE screenshot of this indexing design

```
-----+
| -> Table scan on <temporary> (cost=115.04..117.71 rows=24) (actual time=0.593..0.594 rows=3 loops=1)
|   -> Temporary table with deduplication (cost=114.93..114.93 rows=24) (actual time=0.592..0.592 rows=3 loops=1)
|     -> Nested loop inner join (cost=112.54 rows=24) (actual time=0.468..0.571 rows=3 loops=1)
|       -> Nested loop left join (cost=48.88 rows=24) (actual time=0.406..0.458 rows=5 loops=1)
|         -> Nested loop inner join (cost=40.52 rows=24) (actual time=0.398..0.440 rows=5 loops=1)
|           -> Filter: (Subquery.TimesBagged >= <cache>((0.8 * (select #3)))) (cost=69.72..13.22 rows=78) (actual time=0.379..0.395 rows=6 loops=1)
|             -> Table scan on Subquery (cost=70.47..75.88 rows=234) (actual time=0.156..0.169 rows=102 loops=1)
|               -> Materialize (cost=70.45..70.45 rows=234) (actual time=0.154..0.154 rows=102 loops=1)
|                 -> Group aggregate: count(Bag.ProductId) (cost=47.05 rows=234) (actual time=0.043..0.124 rows=102 loops=1)
|                   -> Covering index scan on Bag using idx_bagitems_productid (cost=23.65 rows=234) (actual time=0.037..0.080 rows=237 loops=1)
|                     -> Select #3 (subquery in condition; run only once)
|                       -> Aggregate: max(Subquery2.TimesBagged) (cost=99.28..99.28 rows=1) (actual time=0.177..0.177 rows=1 loops=1)
|                         -> Table scan on Subquery2 (cost=70.47..75.88 rows=234) (actual time=0.156..0.167 rows=102 loops=1)
|                           -> Materialize (cost=70.45..70.45 rows=234) (actual time=0.155..0.155 rows=102 loops=1)
|                             -> Group aggregate: count(Bag.ProductId) (cost=47.05 rows=234) (actual time=0.043..0.127 rows=102 loops=1)
|                               -> Covering index scan on Bag using idx_bagitems_productid (cost=23.65 rows=234) (actual time=0.041..0.081 rows=237 loops=1)
|                                 -> Filter: ((Pro.LikeCount > 1000) and (Pro.ViewCount > 1)) (cost=0.25 rows=0.3) (actual time=0.007..0.007 rows=1 loops=6)
|                                   -> Single-row index lookup on Pro using PRIMARY (ProductId=Subquery.ProductId) (cost=0.25 rows=1) (actual time=0.006..0.006 rows=1 loops=6)
|                                     -> Single-row index lookup on B using PRIMARY (BrandId=Pro.BrandId) (cost=0.25 rows=1) (actual time=0.003..0.003 rows=1 loops=5)
|                                       -> Limit: 1 row(s) (cost=2.38 rows=1) (actual time=0.022..0.022 rows=1 loops=5)
|                                         -> Filter: (R.Rating > 4) (cost=2.38 rows=3) (actual time=0.022..0.022 rows=1 loops=5)
|                                           -> Index lookup on R using idx_Reviews_ProductId_ReviewId (ProductId=Subquery.ProductId) (cost=2.38 rows=9) (actual time=0.021..0.021 rows=1 loops=5)
|
|-----+
```

- Total Cost: 117.71 + 114.93 + 112.54 + 48.88 + 40.52 + 69.7 + 75.88 + 70.45 + 47.05 + 23.6 + 99.28 + 75.88 + 70.45 + 47.05 + 23.65 + 0.25 + 0.25 + 0.25 + 2.38 + 2.38 + 2.38 = **1045.46**

Pros/Performance Gains: No performance gains.

The LikeCount attribute is only used in one WHERE clause, and so ideally the index should improve filtering performance. However, in this case, LikeCount index is creating performance degradations, perhaps because it is quite large. Other reasons for degradations explained below

```
-----+-----+-----+-----+
| database_name | table_name | index_name | stat_value*@innodb_page_size |
|-----+-----+-----+-----+
| MojoDojoDB   | Products  | idx_Products_LikeCount | 180224 |
|-----+-----+-----+-----+
1 row in set (0.00 sec)
```

Cons/Degradations: Degradations present. Adding this extra index will take up more storage space in the BTREE, and add an extra layer of complexity to queries. Since we already have indexes for all the primary keys, it might not benefit performance significantly to add another index for filtering LikeCount.

Additionally, adding an index can interfere with the MySQL engine developing an efficient query optimization plan. It is possible that because we added this new index, MySQL had to work around it to develop a less efficient query optimization plan, for example, a plan with more locking.

Indexing Design 2

- New Index:
`CREATE INDEX idx_Products_ViewCount ON Products(ViewCount);`
- EXPLAIN ANALYZE screenshot of this indexing design

```

| -> Table scan on <temporary> (cost=53.32..55.21 rows=4) (actual time=0.477..0.477 rows=3 loops=1)
| -> Temporary table with deduplication (cost=52.67..52.67 rows=4) (actual time=0.476..0.476 rows=3 loops=1)
|   -> Nested loop inner join (cost=52.28 rows=4) (actual time=0.369..0.456 rows=3 loops=1)
|     -> Nested loop left join (cost=41.89 rows=4) (actual time=0.318..0.361 rows=3 loops=1)
|       -> Nested loop anti join (cost=40.52 rows=4) (actual time=0.311..0.346 rows=5 loops=1)
|         -> Filter: (Subquery.TimesBagged >= <cache>)(0.8 * (select (3))) (cost=69.72..13.22 rows=78) (actual time=0.299..0.315 rows=6 loops=1)
|           -> Table scan on Subquery (cost=70.47..75.88 rows=234) (actual time=0.143..0.155 rows=102 loops=1)
|             -> Materialize (cost=70.45..70.45 rows=234) (actual time=0.142..0.142 rows=102 loops=1)
|               -> Group aggregate: count(Bag.ProductId) (cost=47.05 rows=234) (actual time=0.031..0.114 rows=102 loops=1)
|                 -> Covering index scan on Bag using idx_bagitems_productid (cost=23.65 rows=234) (actual time=0.026..0.071 rows=237 loops=1)
|                   -> Select #3 (subquery in condition; run only once)
|                     -> Aggregate: max(Subquery2.TimesBagged) (cost=99.28..99.28 rows=1) (actual time=0.139..0.140 rows=1 loops=1)
|                       -> Table scan on Subquery2 (cost=70.47..75.88 rows=234) (actual time=0.119..0.130 rows=102 loops=1)
|                         -> Materialize (cost=70.45..70.45 rows=234) (actual time=0.115..0.118 rows=102 loops=1)
|                           -> Group aggregate: count(Bag.ProductId) (cost=47.05 rows=234) (actual time=0.017..0.095 rows=102 loops=1)
|                             -> Covering index scan on Bag using idx_bagitems_productid (cost=23.65 rows=234) (actual time=0.016..0.056 rows=237 loops=1)
|                               -> Filter: ((Pro.LikeCount > 1000) and (Pro.ViewCount > 0)) (cost=0.25 rows=0.05) (actual time=0.005..0.005 rows=1 loops=6)
|                                 -> Single-row index lookup on Pro using PRIMARY (ProductId=Subquery.ProductId) (cost=0.25 rows=1) (actual time=0.004..0.004 rows=1 loops=6)
|                                   -> Single-row index lookup on Subquery using PRIMARY (BagId=Pro.BagId) (cost=0.28 rows=1) (actual time=0.003..0.003 rows=1 loops=5)
|                                     -> Limit: 1 row(s) (cost=2.45 rows=1) (actual time=0.018..0.018 rows=1 loops=5)
|                                       -> Filter: (R.Rating > 4) (cost=2.45 rows=3) (actual time=0.018..0.018 rows=1 loops=5)
|                                         -> Index lookup on R using idx_Reviews_ProductId_ReviewId (ProductId=Subquery.ProductId) (cost=2.45 rows=9) (actual time=0.018..0.018 rows=1 loops=5)

```

- Total Cost: $52.67 + 55.21 + 52.28 + 41.89 + 40.52 + 69.72 + 75.88 + 70.45 + 47.05 + 23.65 + 99.28 + 75.88 + 70.45 + 47.05 + 23.65 + 0.25 + 0.25 + 0.28 + 2.45 + 2.45 + 2.45 = 853.76$

Pros/Performance Gains: Performance gains present! Since the ViewCount index is used in the main WHERE clause, indexing on ViewCount improves the query's filtering performance.

Cons/Degradations: No cons/degradations. Since we already have a ProductId index in the Brands table (Primary key) and in the Products table (foreign key), it might not benefit performance significantly to add an extra index for ViewCount. However in this case, it does improve performance a slight bit.

Indexing Design 3

- New Index: `CREATE INDEX idx_Reviews_Rating ON Reviews(Rating);`

- EXPLAIN ANALYZE screenshot of this indexing design

```

-----+-----
| -> Table scan on <temporary> (cost=66.95..69.25 rows=9) (actual time=0.599..0.600 rows=3 loops=1)
| -> Temporary table with deduplication (cost=66.65..66.65 rows=9) (actual time=0.598..0.598 rows=3 loops=1)
| -> Nested loop inner join (cost=65.78 rows=9) (actual time=0.467..0.575 rows=3 loops=1)
|   -> Nested loop left join (cost=43.55 rows=9) (actual time=0.401..0.453 rows=5 loops=1)
|     -> Nested loop inner join (cost=40.52 rows=9) (actual time=0.392..0.432 rows=5 loops=1)
|       -> Filter: (Subquery.TimesBagged >= <cache>((0.8 * (select #3)))) (cost=69.72..13.22 rows=78) (actual time=0.371..0.386 rows=6 loops=1)
|         -> Table scan on Subquery (cost=70.47..75.88 rows=234) (actual time=0.194..0.205 rows=102 loops=1)
|           -> Materialize (cost=70.45..70.45 rows=234) (actual time=0.191..0.191 rows=102 loops=1)
|             -> Group aggregate: count(Bag.ProductId) (cost=47.05 rows=234) (actual time=0.036..0.156 rows=102 loops=1)
|               -> Covering index scan on Bag using idx_bagitems_productid (cost=23.65 rows=234) (actual time=0.030..0.105 rows=237 loops=1)
|             -> Select #3 (subquery in condition; run only once) (cost=99.28..99.28 rows=1) (actual time=0.156..0.156 rows=1 loops=1)
|               -> Aggregate: max(Subquery2.TimesBagged) (cost=70.47..75.88 rows=234) (actual time=0.134..0.146 rows=102 loops=1)
|                 -> Materialize (cost=70.45..70.45 rows=234) (actual time=0.134..0.134 rows=102 loops=1)
|                   -> Group aggregate: count(Bag.ProductId) (cost=47.05 rows=234) (actual time=0.020..0.105 rows=102 loops=1)
|                     -> Covering index scan on Bag using idx_bagitems_productid (cost=23.65 rows=234) (actual time=0.019..0.061 rows=237 loops=1)
|                   -> Filter: ((Pro.LikeCount > 1000) and (Pro.ViewCount > 1)) (cost=0.25 rows=0.1) (actual time=0.007..0.007 rows=1 loops=6)
|                     -> Single-row index lookup on Pro using PRIMARY (ProductId=Subquery.ProductId) (cost=0.25 rows=1) (actual time=0.006..0.007 rows=1 loops=6)
|                   -> Single-row index lookup on B using PRIMARY (BrandId=Pro.BrandId) (cost=0.26 rows=1) (actual time=0.004..0.004 rows=1 loops=5)
|                   -> Limit: 1 row(s) (cost=2.42 rows=1) (actual time=0.024..0.024 rows=1 loops=5)
|                   -> Filter: (R.Rating > 4) (cost=2.42 rows=5) (actual time=0.024..0.024 rows=1 loops=5)
|                   -> Index lookup on R using idx_Reviews_ProductId_ReviewId (ProductId=Subquery.ProductId) (cost=2.42 rows=9) (actual time=0.023..0.023 rows=1 loops=5)

```

- Total Cost: 69.25 + 66.65 + 65.78 + 43.55 + 40.52 + 69.72 + 75.88 + 70.45 + 47.0 + 23.65 + 99.28 + 75.88 + 70.45 + 47.05 + 23.65 + 0.25 + 0.25 + 0.26 + 2.42 + 2.42 + 2.42 = **896.78**

- Pros and cons, performance gains and degradations of this indexing design

Pros/Performance Gains: Performance gains present! Since the Rating index is used in the main WHERE clause, indexing on Rating improves the query's filtering performance.

Cons/Degradations: No cons/degradations. Since we already have a ProductId index in the Reviews table (Primary key) and in the Products table (foreign key), it might not benefit performance significantly to add an extra index for Rating. But, because we are rigorously filtering on rating, it improves performance.

Final Indexing Choice and Why

Indexing Design #2 is our final indexing choice as it has the lowest cost — 853 as opposed to the original cost of 906 without the index. Because our query includes ViewCount in the WHERE clause, this query will filter thousands of records by ViewCount. Adding an index on this non-primary key attribute will improve the performance of this filtering

Query 2: User Bag Overview During Friends Search

```
SELECT UserName, COUNT(ProductId) AS NumProductsInBag
FROM Users LEFT OUTER JOIN BagItems ON Users.UserId = BagItems.UserId
GROUP BY (Users.UserId)
ORDER BY LEAST(LEVENSHTEIN(UserName, *search*), LEVENSHTEIN(FirstName,
*search*), LEVENSHTEIN(LastName, *search*))
LIMIT 15;
```

Output: example query replaces *search* with “kylie”

UserName	NumProductsInBag
kyliejenner	4
julie	3
karlie	4
ria	4
natalie	4
bellahadid	4
elvis	3
Hyrain	4
tati	1
i.love.skincare	4
taylor	2
nitya	3
oju	4
marilyn	4
i.love.makeup	4

15 rows in set (0.02 sec)

EXPLAIN ANALYZE Performance before adding indexes:

```
| -> Limit: 15 row(s) (actual time=15.981..15.983 rows=15 loops=1)
|   -> Sort: least(LEVENSHTEIN(Users.UserName, 'kylie'), LEVENSHTEIN(Users.FirstName, 'kylie'), LEVENSHTEIN(Users.LastName, 'kylie')), limit input to 15 row(s) per chunk (actual time=15.980..15.981 rows=15 loops=1)
|     -> Table scan on <temporary> (actual time=15.957..15.961 rows=21 loops=1)
|       -> Aggregate using temporary table (actual time=15.955..15.955 rows=21 loops=1)
|         -> Nested loop left join (cost=15.52 rows=73) (actual time=0.079..0.286 rows=73 loops=1)
|           -> Table scan on Users (cost=2.35 rows=21) (actual time=0.040..0.063 rows=21 loops=1)
|             -> Covering index lookup on BagItems using PRIMARY (UserId=Users.UserId) (cost=0.30 rows=3) (actual time=0.005..0.009 rows=3 loops=21)
```

Total Cost: 15.52 + 2.35 + 0.30 = 18.17

Indexing Design 1

- New Index: CREATE INDEX idx_Users_UserName ON Users (UserName);
- EXPLAIN ANALYZE screenshot of this indexing design

```
| -> Limit: 15 row(s) (actual time=14.969..14.971 rows=15 loops=1)
|   -> Sort: least(LEVENSHTEIN(Users.UserName, 'kylie'), LEVENSHTEIN(Users.FirstName, 'kylie'), LEVENSHTEIN(Users.LastName, 'kylie')), limit input to 15 row(s) per chunk (actual time=14.968..14.969 rows=15 loops=1)
|     -> Table scan on <temporary> (actual time=14.936..14.940 rows=21 loops=1)
|       -> Aggregate using temporary table (actual time=14.934..14.934 rows=21 loops=1)
|         -> Nested loop left join (cost=15.52 rows=73) (actual time=0.039..0.196 rows=73 loops=1)
|           -> Table scan on Users (cost=2.35 rows=21) (actual time=0.025..0.041 rows=21 loops=1)
|             -> Covering index lookup on BagItems using PRIMARY (UserId=Users.UserId) (cost=0.30 rows=3) (actual time=0.004..0.007 rows=3 loops=21)
```

- Total Cost: 15.52 + 2.35 + 0.30 = 18.17
- Pros and cons, performance gains and degradations of this indexing design

We expect no performance improvement or degradation. This is likely due to the fact that Users.UserName is only used in the Levenshtein distance calculation. Due to the primary key

UserId in the User table, and the fact that the UserId is the only attribute being used in the JOIN and GROUP BY clauses, adding an index on the UserName attribute of the Users table will only take up more storage space without improving performance.

Indexing Design 2

- New Index: CREATE INDEX idx_Users_CoveringIndex ON Users (UserName, FirstName, LastName);
- EXPLAIN ANALYZE screenshot of this indexing design

```
| -> Limit: 15 row(s) (actual time=15.724..15.726 rows=15 loops=1)
| -> Sort: least (LEVENSHTEIN(Users.UserName,'kylie'),LEVENSHTEIN(Users.FirstName,'kylie'),LEVENSHTEIN(Users.LastName,'kylie')), limit input to 15 row(s) per chunk (actual time=15.723..15.724 rows=15 loops=1)
| -> Table scan on <temporary> (actual time=15.699..15.704 rows=21 loops=1)
| -> Aggregate using temporary table (actual time=15.697..15.697 rows=21 loops=1)
| -> Nested loop left join (cost=15.52 rows=73) (actual time=0.038..0.220 rows=73 loops=1)
| ->   Covering index scan on Users using idx_Users_CoveringIndex (cost=2.35 rows=21) (actual time=0.027..0.043 rows=21 loops=1)
| ->   Covering index lookup on BagItems using PRIMARY (UserId=Users.UserId) (cost=0.30 rows=3) (actual time=0.005..0.008 rows=3 loops=21)
```

- Total Cost: 15.52 + 2.35 + 0.30 = 18.17
- Pros and cons, performance gains and degradations of this indexing design

We expect no performance improvement or degradation. This is likely due to the fact that Users.UserName, Users.FirstName, and Users.LastName are only used in the Levenshtein distance calculations. Due to the primary key UserId in the User table, and the fact that the UserId is the only attribute being used in the JOIN and GROUP BY clauses, adding an index on the UserName, FirstName, and LastName attributes of the Users table will only take up more storage space without improving performance.

Indexing Design 3

- New Index: CREATE INDEX idx_Users_BagItems_JoinIndex ON BagItems (UserId, ProductId);
- EXPLAIN ANALYZE screenshot of this indexing design

```
| -> Limit: 15 row(s) (actual time=14.999..15.001 rows=15 loops=1)
| -> Sort: least (LEVENSHTEIN(Users.UserName,'kylie'),LEVENSHTEIN(Users.FirstName,'kylie'),LEVENSHTEIN(Users.LastName,'kylie')), limit input to 15 row(s) per chunk (actual time=14.998..14.999 rows=15 loops=1)
| -> Table scan on <temporary> (actual time=14.967..14.973 rows=21 loops=1)
| -> Aggregate using temporary table (actual time=14.964..14.964 rows=21 loops=1)
| -> Nested loop left join (cost=15.52 rows=73) (actual time=0.050..0.215 rows=73 loops=1)
| ->   Table scan on Users (cost=2.35 rows=21) (actual time=0.035..0.055 rows=21 loops=1)
| ->   Covering index lookup on BagItems using PRIMARY (UserId=Users.UserId) (cost=0.30 rows=3) (actual time=0.004..0.007 rows=3 loops=21)
```

- Total Cost: 15.52 + 2.35 + 0.30 = 18.17
- Pros and cons, performance gains and degradations of this indexing design

We expect no performance improvement or degradation. This is likely due to the fact that BagItems.UserId is used in the JOIN clause but is already a primary key, and BagItems.ProductId is not used in the query. Since ProductId is not used in the query, adding an index on it only takes up more storage space without improving performance.

Final Indexing Choice and Why

Justification: The final index design we chose for Advanced Query #2 is the default index (the primary key UserId in the Users table and the primary/foreign key UserId in the BagItems table). We chose this index because all of the other indexing designs we tried did not impact the cost of the query. This is because this query simply uses the primary key UserId in both the JOIN and GROUP BY clauses, while any other attributes are simply used in the SELECT or ORDER BY clauses, and thus do not have an effect on the cost.

Query 3: Users also bagged

```
SELECT BI.ProductId AS OtherPID,
       P.ProductName,
       AVG(R.Rating) AS AverageRating,
       COUNT(R.ReviewId) AS NumberOfReviews
FROM BagItems BI
     JOIN Products P ON BI.ProductId = P.ProductId
     JOIN
     (
         SELECT ProductId, MAX(ReviewId) AS LatestReviewId
         FROM Reviews
         GROUP BY ProductId
     ) AS LatestReview
     ON BI.ProductId = LatestReview.ProductId
     JOIN Reviews R
     ON LatestReview.LatestReviewId = R.ReviewId
WHERE BI.ProductId != *current-product*
     AND BI.UserId IN (
         SELECT UserId FROM BagItems WHERE ProductId = *current-product*)
GROUP BY BI.ProductId, P.ProductName
ORDER BY BI.ProductId;
```

```
SELECT BI.ProductId AS OtherPID, P.ProductName, AVG(R.Rating) AS
AverageRating, COUNT(R.ReviewId) AS NumberOfReviews FROM BagItems BI
JOIN Products P ON BI.ProductId = P.ProductId JOIN (SELECT
ProductId, MAX(ReviewId) AS LatestReviewId FROM Reviews GROUP BY
ProductId) AS LatestReview ON BI.ProductId = LatestReview.ProductId
JOIN Reviews R ON LatestReview.LatestReviewId = R.ReviewId WHERE
BI.ProductId != *current-product* AND BI.UserId IN (SELECT UserId FROM
BagItems WHERE ProductId = *current-product*) GROUP BY BI.ProductId,
P.ProductName ORDER BY BI.ProductId;
```

Output: example query replaces *current-product* with 'P501265'

OtherPID	ProductName	AverageRating	NumberOfReviews
P392235	The Camellia Oil 2-in-1 Makeup Remover & Cleanser	5.0000	15
P429659	Squalane + Hyaluronic Toning Mist	5.0000	1
P438643	The Balance pH Balancing Gel Cleanser	5.0000	1
P441644	Mini Superfood Antioxidant Cleanser	2.0000	17
P465741	Wild Huckleberry 8-Acid Polishing Peel Mask	5.0000	1
P480278	Rapid Radiance Set	4.0000	1
P481084	Mini Revitalizing Supreme+ Youth Power Creme Moisturizer	5.0000	17
P481817	Beauty Elixir Prep, Set, Glow Face Mist	4.0000	1
P505020	The POREfessional Good Cleanup Foaming Cleanser	5.0000	1

9 rows in set (0.03 sec)

EXPLAIN ANALYZE Performance before adding indexes:

Total Cost: 1041.91 + 554.88 + 67.85 + 21.95 + 2.14 + 0.38 + 0.38 + 0.25 + 826.25 + 911.95 + 731.85 + 0.25 = 4160.04

```
| -> Sort: BI.ProductId (actual time=16.188..16.189 rows=9 loops=1)
|   -> Table scan on <temporary> (actual time=16.167..16.169 rows=9 loops=1)
|     -> Aggregate using temporary table (actual time=16.165..16.165 rows=9 loops=1)
|       -> Nested loop inner join (cost=1041.91 rows=1392) (actual time=15.947..16.085 rows=55 loops=1)
|         -> Nested loop inner join (cost=554.88 rows=1392) (actual time=15.334..15.977 rows=55 loops=1)
|           -> Nested loop inner join (cost=67.85 rows=131) (actual time=0.033..0.409 rows=146 loops=1)
|             -> Nested loop inner join (cost=21.95 rows=131) (actual time=0.026..0.161 rows=146 loops=1)
|               -> Covering index lookup on BagItems using idx_bagitems_productid (ProductId='P501265') (cost=2.14 rows=17) (actual time=0.015..0.022 rows=17 loops=1)
|               -> Filter: (BI.ProductId <> 'P501265') (cost=0.38 rows=8) (actual time=0.004..0.007 rows=9 loops=17)
|               -> Covering index lookup on BI using PRIMARY (UserId=BagItems.UserId) (cost=0.38 rows=8) (actual time=0.004..0.006 rows=10 loops=17)
|             -> Single-row index lookup on P using PRIMARY (ProductId=BI.ProductId) (cost=0.25 rows=1) (actual time=0.001..0.002 rows=1 loops=146)
|           -> Filter: (LatestReview.LatestReviewId is not null) (cost=826.25..2.66 rows=11) (actual time=0.106..0.106 rows=0 loops=146)
|           -> Index lookup on LatestReview using <auto_key> (ProductId=BI.ProductId) (actual time=0.106..0.106 rows=0 loops=146)
|           -> Materialize (cost=911.95..911.95 rows=1801) (actual time=15.271..15.271 rows=1800 loops=1)
|           -> Covering index skip scan for grouping on Reviews using ProductId (cost=731.85 rows=1801) (actual time=0.018..12.924 rows=1800 loops=1)
|         -> Single-row index lookup on R using PRIMARY (ReviewId=LatestReview.LatestReviewId) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=55)
```

Indexing Design 1

- New Index: CREATE INDEX ByRating on Reviews(Rating);
- EXPLAIN ANALYZE screenshot of this indexing design

```
| -> Sort: BI.ProductId (actual time=27.653..27.654 rows=9 loops=1)
|   -> Table scan on <temporary> (actual time=27.622..27.625 rows=9 loops=1)
|     -> Aggregate using temporary table (actual time=27.620..27.620 rows=9 loops=1)
|       -> Nested loop inner join (cost=1041.91 rows=1392) (actual time=25.505..27.501 rows=55 loops=1)
|         -> Nested loop inner join (cost=554.88 rows=1392) (actual time=25.489..27.387 rows=55 loops=1)
|           -> Nested loop inner join (cost=67.85 rows=131) (actual time=0.122..1.638 rows=146 loops=1)
|             -> Nested loop inner join (cost=21.95 rows=131) (actual time=0.112..0.538 rows=146 loops=1)
|               -> Covering index lookup on BagItems using idx_bagitems_productid (ProductId='P501265') (cost=2.14 rows=17) (actual time=0.021..0.030 rows=17 loops=1)
|               -> Filter: (BI.ProductId <> 'P501265') (cost=0.38 rows=8) (actual time=0.025..0.029 rows=9 loops=17)
|               -> Covering index lookup on BI using PRIMARY (UserId=BagItems.UserId) (cost=0.38 rows=8) (actual time=0.024..0.027 rows=10 loops=17)
|             -> Single-row index lookup on P using PRIMARY (ProductId=BI.ProductId) (cost=0.25 rows=1) (actual time=0.007..0.007 rows=1 loops=146)
|           -> Filter: (LatestReview.LatestReviewId is not null) (cost=826.25..2.66 rows=11) (actual time=0.176..0.176 rows=0 loops=146)
|           -> Index lookup on LatestReview using <auto_key> (ProductId=BI.ProductId) (actual time=0.176..0.176 rows=0 loops=146)
|           -> Materialize (cost=911.95..911.95 rows=1801) (actual time=25.324..25.324 rows=1800 loops=1)
|           -> Covering index skip scan for grouping on Reviews using ProductId (cost=731.85 rows=1801) (actual time=0.154..22.628 rows=1800 loops=1)
|         -> Single-row index lookup on R using PRIMARY (ReviewId=LatestReview.LatestReviewId) (cost=0.25 rows=1) (actual time=0.002..0.002 rows=1 loops=55)
```

- Total Cost: 1041.91 + 554.88 + 67.85 + 21.95 + 2.14 + 0.38 + 0.38 + 0.25 + 826.25 + 911.95 + 731.85 + 0.25 = 4160.04
- Pros and cons, performance gains and degradations of this indexing design:

This indexing design is not ideal because it does not significantly enhance the performance for queries where Rating is not used directly for filtering, sorting, or as a crucial part of join conditions, despite its presence in the query output. Although the rating is outputted in the query, it is not directly involved in filtering or sorting and it is not a part of major JOIN operations. Although the output showed no performance gain or degradation, this index has the potential to slow other operations because it would need to be updated based on modifications to the Reviews table. Since this index does not lead to performance gains, it could take up unnecessary storage and slow other operations.

Indexing Design 2

- New Index: CREATE INDEX ByName on Products(ProductName);
- EXPLAIN ANALYZE screenshot of this indexing design

```
| -> Sort: BI.ProductId (actual time=16.714..16.715 rows=9 loops=1)
|   -> Table scan on <temporary> (actual time=16.693..16.694 rows=9 loops=1)
|     -> Aggregate using temporary table (actual time=16.691..16.691 rows=9 loops=1)
|       -> Nested loop inner join (cost=1041.91 rows=1392) (actual time=15.874..16.611 rows=55 loops=1)
|         -> Nested loop inner join (cost=554.88 rows=1392) (actual time=15.859..16.517 rows=55 loops=1)
|           -> Nested loop inner join (cost=67.85 rows=131) (actual time=0.034..0.403 rows=146 loops=1)
|             -> Nested loop inner join (cost=21.95 rows=131) (actual time=0.026..0.158 rows=146 loops=1)
|               -> Covering index lookup on BagItems using idx_bagitems_productid (ProductId='P501265') (cost=2.14 rows=17) (actual time=0.015..0.021 rows=17 loops=1)
|               -> Filter: (BI.ProductId <> 'P501265') (cost=0.38 rows=8) (actual time=0.004..0.007 rows=9 loops=17)
|               -> Covering index lookup on BI using PRIMARY (UserId=BagItems.UserId) (cost=0.38 rows=8) (actual time=0.004..0.006 rows=10 loops=17)
|             -> Single-row index lookup on P using PRIMARY (ProductId=BI.ProductId) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=146)
|           -> Filter: (LatestReview.LatestReviewId is not null) (cost=826.25..2.66 rows=11) (actual time=0.110..0.110 rows=0 loops=146)
|           -> Index lookup on LatestReview using <auto_key> (ProductId=BI.ProductId) (actual time=0.110..0.110 rows=0 loops=146)
|           -> Materialize (cost=911.95..911.95 rows=1801) (actual time=15.793..15.793 rows=1800 loops=1)
|           -> Covering index skip scan for grouping on Reviews using ProductId (cost=731.85 rows=1801) (actual time=0.018..13.367 rows=1800 loops=1)
|         -> Single-row index lookup on R using PRIMARY (ReviewId=LatestReview.LatestReviewId) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=55)
```

- Total Cost: $1041.91 + 554.88 + 67.85 + 21.95 + 2.14 + 0.38 + 0.38 + 0.25 + 826.25 + 911.95 + 731.85 + 0.25 = 4160.04$
- Pros and cons, performance gains and degradations of this indexing design: This indexing design did not facilitate any performance gains or degradations. The impact of the index remains neutral after the analysis, summarized with a total cost of 4160.04, underscores that this particular indexing design neither improves nor detracts from the query's efficiency. Since ProductName is often referenced in the scope of our application, it could be useful for other queries, but this query does not involve ProductName in major operations. The query identifies distinct products based on ProductId, not based on ProductName, so the index does not affect the cost.

Indexing Design 3

- New Index: CREATE INDEX ByProduct on Products(ProductId, ProductName);
- EXPLAIN ANALYZE screenshot of this indexing design

```

| -> Sort: BI.ProductId (actual time=16.054..16.054 rows=9 loops=1)
    -> Table scan on <temporaz> (actual time=16.031..16.033 rows=9 loops=1)
        -> Aggregate using temporary table (actual time=16.029..16.029 rows=9 loops=1)
            -> Nested loop inner join (cost=1041.91 rows=1392) (actual time=15.245..15.951 rows=55 loops=1)
                -> Nested loop inner join (cost=554.88 rows=1392) (actual time=15.231..15.857 rows=55 loops=1)
                    -> Nested loop inner join (cost=67.85 rows=131) (actual time=0.033..0.401 rows=146 loops=1)
                        -> Nested loop inner join (cost=21.95 rows=131) (actual time=0.022..0.155 rows=146 loops=1)
                            -> Covering index lookup on BagItems using idx_bagitems_productid (ProductId='P501265') (cost=2.14 rows=17) (actual time=0.015..0.021 rows=17 loops=1)
                                -> Filter: (BI.ProductId <> 'P501265') (cost=0.38 rows=8) (actual time=0.004..0.007 rows=9 loops=17)
                                    -> Covering index lookup on BI using PRIMARY (UserId=BagItems.UserId) (cost=0.38 rows=8) (actual time=0.004..0.006 rows=10 loops=17)
                                        -> Single-row index lookup on P using PRIMARY (ProductId=BI.ProductId) (cost=0.25 rows=1) (actual time=0.001..0.002 rows=1 loops=146)
                                            -> Filter: (LatestReview.latestReviewId is not null) (cost=926.25..2.46 rows=11) (actual time=0.106..0.106 rows=0 loops=146)
                                                -> Index lookup on LatestReview using kauto_Key1> (ProductId=BI.ProductId) (actual time=0.105..0.105 rows=0 loops=146)
                                                    -> Materialize (cost=911.95..911.95 rows=1801) (actual time=15.165..15.165 rows=1800 loops=1)
                                                        -> Covering index skip scan for grouping on Reviews using ProductId (cost=731.85 rows=1801) (actual time=0.019..12.885 rows=1800 loops=1)
                                                            -> Single-row index lookup on R using PRIMARY (ReviewId=LatestReview.LatestReviewId) (cost=0.25 rows=1) (actual time=0.001..0.001 rows=1 loops=55)

```

- Total Cost: $1041.91 + 554.88 + 67.85 + 21.95 + 2.14 + 0.38 + 0.38 + 0.25 + 826.25 + 911.95 + 731.85 + 0.25 = 4160.04$
- Pros and cons, performance gains and degradations of this indexing design: There was no performance gain or degradation due to the creation of this index. Although the output depicted no change in cost, this indexing design has the most potential to induce performance improvement. Since ProductId is a part of this index, the database will utilize this index for the JOIN operation. One con is that indexing with multiple columns consumes more storage space and manages more data, so applying a modification may be more costly. Specifically, due to the inclusion of two columns, updates to either ProductId or ProductName would necessitate index maintenance.

Final Indexing Choice and Why

- Report on the final index design you selected and explain why you chose it, referencing the analysis you performed:

After thorough analysis, none of my indexing designs (ByRating, ByName, and ByProduct) demonstrated any performance improvement for the query. As a result, I would not select any of the indexing designs because they would take up more storage space rather than saving cost. Specifically, the indexes on Rating and ProductName did not align with the primary attributes used in critical operations such as joins and filtering. Due to this, the indices did not impact the cost. Although the last indexing design had one column that indexed on ProductId, this still did not display any explicit performance gain or degradation, so I will not select this indexing design.

Query 4: Aggregations for common products bagged between the application user and the user's bag they are viewing

```
SELECT BI1.UserId AS UserId1, BI2.UserId AS UserId2,
COUNT(BI1.ProductId) AS SharedProductsCount, GROUP_CONCAT(DISTINCT
P.ProductName ORDER BY P.ProductName SEPARATOR ', ') AS
SharedProductNames, AVG(P.Price) AS AveragePriceOfSharedProducts
FROM BagItems BI1
JOIN BagItems BI2 ON BI1.ProductId = BI2.ProductId AND BI1.UserId <
BI2.UserId
JOIN Products P ON BI1.ProductId = P.ProductId
GROUP BY BI1.UserId, BI2.UserId
HAVING COUNT(BI1.ProductId) > 0
LIMIT 15;
```

Output:

UserId1	UserId2	SharedProductsCount	SharedProductNames	AveragePriceOfSharedProducts
1	2	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	3	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	4	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	5	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	6	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	7	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	8	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	9	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	11	5	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	36.6000
ascara, Soft Glam Eyeshadow Palette				
1	13	5	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	36.6000
ascara, Soft Glam Eyeshadow Palette				
1	16	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	17	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	18	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	19	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			
1	20	6	Lash Extension Tubing Mascara, Mini Revitalizing Supreme+ Youth Power Creme	38.8333
ascara, Soft Glam Eyeshadow Palette,	The Camellia Oil 2-in-1 Makeup Remover & Cleanser			

EXPLAIN ANALYZE Performance before adding indexes:

```
--+
| -> Limit: 15 row(s) (actual time=2.527..2.578 rows=15 loops=1)
|   -> Filter: (count(BagItems.ProductId) > 0) (actual time=2.527..2.577 rows=15 loops=1)
|     -> Group aggregate: count(BagItems.ProductId), count(BagItems.ProductId), group_concat(distinct Products.ProductName order by Products.ProductName ASC separator ',
products.Price) (actual time=2.526..2.574 rows=15 loops=1)
|       -> Sort: BI1.UserId, BI2.UserId (actual time=2.503..2.512 rows=89 loops=1)
|         -> Stream results (cost=153.87 rows=138) (actual time=0.076..2.280 rows=861 loops=1)
|           -> Nested loop inner join (cost=153.87 rows=138) (actual time=0.074..1.790 rows=861 loops=1)
|             -> Nested loop inner join (cost=105.49 rows=138) (actual time=0.064..1.384 rows=861 loops=1)
|               -> Covering index scan on BI1 using ProductId (cost=17.65 rows=174) (actual time=0.041..0.072 rows=179 loops=1)
|               -> Filter: (BI1.UserId < BI2.UserId) (cost=0.27 rows=1) (actual time=0.004..0.007 rows=5 loops=179)
|                 -> Covering index lookup on BI2 using ProductId (ProductId=BI1.ProductId) (cost=0.27 rows=2) (actual time=0.002..0.006 rows=11 loops=179)
|                   -> Single-row index lookup on P using PRIMARY (ProductId=BI1.ProductId) (cost=0.25 rows=1) (actual time=0.000..0.000 rows=1 loops=861)
|
```

Total Cost: 153.87+153.87+105.49+17.65+0.27+0.27 + 0.25= **431.67**

Indexing Design 1

- New Index: CREATE INDEX idx_bagitems_productid ON BagItems(ProductId);
- EXPLAIN ANALYZE screenshot of this indexing design

```
| -> Filter: (count(BagItems.ProductId) > 0) (actual time=3.444..3.979 rows=171 loops=1)
|   -> Group aggregate: count(BagItems.ProductId), count(BagItems.ProductId), group_concat(distinct Products.ProductName order by Products.ProductName ASC separator ' ', '), avg(Products.Price) (actual time=3.442..3.962 rows=171 loops=1)
|     -> Sort: B11.UserId, B12.UserId (actual time=3.422..3.507 rows=861 loops=1)
|       -> Stream results (cost=153.87 rows=138) (actual time=0.210..3.177 rows=861 loops=1)
|         -> Nested loop inner join (cost=153.87 rows=138) (actual time=0.207..2.689 rows=861 loops=1)
|           -> Nested loop inner join (cost=105.49 rows=138) (actual time=0.194..2.284 rows=861 loops=1)
|             -> Covering index scan on B11 using idx_bagitems_productid (cost=17.65 rows=174) (actual time=0.163..0.207 rows=179 loops=1)
|               -> Filter: (B11.UserId < B12.UserId) (cost=0.27 rows=1) (actual time=0.009..0.011 rows=5 loops=179)
|                 -> Covering index lookup on B12 using idx_bagitems_productid (ProductId=B11.ProductId) (cost=0.27 rows=2) (actual time=0.006..0.010 rows=11 loops=179)
|                   -> Single-row index lookup on P using PRIMARY (ProductId=B11.ProductId) (cost=0.25 rows=1) (actual time=0.000..0.000 rows=1 loops=861)
```

- Total Cost: 153.87+153.87+105.49+17.65+0.27+0.27 + 0.25= **431.67**
- Pros and cons, performance gains and degradations of this indexing design
 - There are no performance improvements or degradations for this index. This is most likely because ProductId is both a primary key and a foreign key to another table, so it is possible that, being a primary key, it already provided sufficient indexing for the query's needs without the index. Being that ProductId is also a foreign key in BagItems, this index is now tied to a foreign key constraint and will not be removed. The indexing design is used in the query's execution plan, and does not have any negative effects on the performance of the query.

Indexing Design 2

- New Index: CREATE INDEX idx_products_productname ON Products(ProductName);
- EXPLAIN ANALYZE screenshot of this indexing design

```
-----
| -> Filter: (count(BagItems.ProductId) > 0) (actual time=2.702..3.273 rows=171 loops=1)
|   -> Group aggregate: count(BagItems.ProductId), count(BagItems.ProductId), group_concat(distinct Products.ProductName order by Products.ProductName ASC separator ' ', '), avg(Products.Price) (actual time=2.701..3.256 rows=171 loops=1)
|     -> Sort: B11.UserId, B12.UserId (actual time=2.684..2.790 rows=861 loops=1)
|       -> Stream results (cost=153.87 rows=138) (actual time=0.055..2.426 rows=861 loops=1)
|         -> Nested loop inner join (cost=153.87 rows=138) (actual time=0.053..1.943 rows=861 loops=1)
|           -> Nested loop inner join (cost=105.49 rows=138) (actual time=0.045..1.500 rows=861 loops=1)
|             -> Covering index scan on B11 using idx_bagitems_productid (cost=17.65 rows=174) (actual time=0.027..0.063 rows=179 loops=1)
|               -> Filter: (B11.UserId < B12.UserId) (cost=0.27 rows=1) (actual time=0.005..0.008 rows=5 loops=179)
|                 -> Covering index lookup on B12 using idx_bagitems_productid (ProductId=B11.ProductId) (cost=0.27 rows=2) (actual time=0.002..0.007 rows=11 loops=179)
|                   -> Single-row index lookup on P using PRIMARY (ProductId=B11.ProductId) (cost=0.25 rows=1) (actual time=0.000..0.000 rows=1 loops=861)
```

- Total Cost: 153.87+153.87+105.49+17.65+0.27+0.27 + 0.25= **431.67**
- Pros and cons, performance gains and degradations of this indexing design
 - There are no performance improvements or degradations for this index. It can be seen above that the query's execution plan doesn't use the index at all. This could be because the computational cost is heavily dominated by something other than retrieving or sorting by ProductName, so the impact of indexing on ProductName is negligible and the ProductName isn't much of a limiting factor.

Indexing Design 3

- New Index: CREATE INDEX idx_products_productname_price ON Products(ProductName, Price);
- EXPLAIN ANALYZE screenshot of this indexing design

```
-----
| -> Filter: (count(BagItems.ProductId) > 0) (actual time=2.744..3.310 rows=171 loops=1)
|   -> Group aggregate: count(BagItems.ProductId), count(BagItems.ProductId), group_concat(distinct Products.ProductName order by Products.ProductName ASC separator ' ', '), avg(Products.Price) (actual time=2.743..3.292 rows=171 loops=1)
|     -> Sort: B11.UserId, B12.UserId (actual time=2.726..2.843 rows=861 loops=1)
|       -> Stream results (cost=153.87 rows=138) (actual time=0.061..2.476 rows=861 loops=1)
|         -> Nested loop inner join (cost=153.87 rows=138) (actual time=0.059..1.967 rows=861 loops=1)
|           -> Nested loop inner join (cost=105.49 rows=138) (actual time=0.051..1.538 rows=861 loops=1)
|             -> Covering index scan on B11 using idx_bagitems_productid (cost=17.65 rows=174) (actual time=0.032..0.077 rows=179 loops=1)
|               -> Filter: (B11.UserId < B12.UserId) (cost=0.27 rows=1) (actual time=0.005..0.008 rows=5 loops=179)
|                 -> Covering index lookup on B12 using idx_bagitems_productid (ProductId=B11.ProductId) (cost=0.27 rows=2) (actual time=0.002..0.007 rows=11 loops=179)
|                   -> Single-row index lookup on P using PRIMARY (ProductId=B11.ProductId) (cost=0.25 rows=1) (actual time=0.000..0.000 rows=1 loops=861)
```

- Total Cost: 153.87+153.87+105.49+17.65+0.27+0.27 + 0.25= **431.67**
- Pros and cons, performance gains and degradations of this indexing design

- There are no performance gains and degradations for this indexing design. The query's execution plan doesn't use this indexing design either. Indexing ProductName and Price together did not outperform the current execution plan significantly enough, which could mean the cost of using the index may have outweighed the performance benefits.

Final Indexing Choice and Why

- Report on the final index design you selected and explain why you chose it, referencing the analysis you performed
 - In the final index design, we will be using Indexing Design 1 and not using Indexing Design 2 and 3. Although Indexing Design 1 doesn't change the computational cost, the index cannot be removed without changing the entire database due to a Foreign Key constraint. It is utilized in the query's execution plan, therefore we decided to continue using it. It also does not affect the total cost negatively. We will not be using index designs 2 and 3 because the impact of them is clearly negligible, and using the index designs will not outperform a full table scan when executing the query. This is likely due to another part of the query taking up the majority of the computational cost. Using the GROUP_CONCAT function in this query was likely very costly and time consuming compared to the indexes attempted, however there is little indexing that can be done to lower the cost of the GROUP_CONCAT usage.

Appendix

QUERIES IN PROGRESS/OLD QUERIES

#	Advanced query	Screenshot of max(15, query result) rows of advanced query results
1	Average Rating // SELECT "P132239" FROM SELECT ProductId, AVG(Rating), COUNT(Rating) AS NumRatings FROM Products Pro JOIN Reviews	On the Products page, to display these stats ★★★★★ 4.97 avg rating (from 19 reviews)

	Rev ON Pro.ProductId = Rev.ProductId GROUP BY ProductId;	
4a	Seeing how many items from each brand a user has in their bag: SELECT U.UserId, U.FirstName, U.LastName, COUNT(BI.ProductId) AS ProductCount, P.BrandId FROM Users U JOIN BagItems BI ON U.UserId = BI.UserId JOIN Products P ON BI.ProductId = P.ProductId GROUP BY U.UserId, P.BrandId HAVING COUNT(BI.ProductId) > 3 LIMIT 15;	
4b *	Comparing one user's bag to another user's to see how many items they have similar: SELECT BI1.UserId AS User1, BI2.UserId AS User2, COUNT(BI1.ProductId) AS SharedProductCount FROM BagItems BI1 JOIN BagItems BI2 ON BI1.ProductId = BI2.ProductId AND BI1.UserId <> BI2.UserId WHERE BI1.UserId = 18 AND BI2.UserId = 6 GROUP BY BI1.UserId, BI2.UserId HAVING COUNT(BI1.ProductId) > 2 LIMIT 15;	